



Shared-profile CUDA validation: execution works, speedup not yet observed

Latency



Throughput



RTX 3090 · 8 requests · max batch size 4 · 1 warmup + 3 repeats · effective bits = 6.0 for both policies.